ROADMAP v2.0

Databricks Certified

Data Engineer Associate

Preparation Roadmap · Based on May 2026 Exam Guide

45 Questions · 90 Minutes · USD 200 · Valid 2 Years

■ Exam Snapshot

Item	Detail
Questions	45 scored multiple-choice
Duration	90 minutes
Passing Score	~70% (industry standard)
Cost	USD 200
Delivery	Online or test center
Prerequisites	None (6 months Databricks experience recommended)
Validity	2 years (re-certify with current exam)

■ Exam Domains & Weights (May 2026)

#	Domain	Weight	Focus Area
1	Databricks Intelligence Platform	6%	Architecture, Delta Lake, Unity Catalog, Compute
2	Data Ingestion & Loading	21%	COPY INTO, Auto Loader, Lakeflow Connect
3	Data Transformation & Modeling	22%	PySpark/SQL ETL, joins, Bronze→Silver→Gold
4	Working with Lakeflow Jobs	16%	Orchestration, control flow, triggers, DAG
5	Implementing CI/CD	10%	Git Folders, Automation Bundles (DABs), CLI
6	Troubleshooting & Optimization	10%	Spark UI, skew, OOM, Liquid Clustering
7	Governance & Security	15%	Unity Catalog, GRANT/REVOKE, masking, ABAC

■ **Heaviest sections: Data Transformation (22%)** + **Data Ingestion (21%)** = 43% of exam. Combined with Governance (15%), these three domains cover **58%** of the test.

ROADMAP v2.0

Databricks Certified

45 Questions · 90 Minutes · USD 200 · Valid 2 Years

PHASE 1 · Weeks 1–2	Platform Foundation (Domain 1)	6%
Wk 1	Databricks Data Intelligence Platform overview · Lakehouse architecture · Workspace UI · Delta Lake fundamentals (ACID, time travel)	
Wk 2	Unity Catalog basics: metastore, catalog, schema, tables · Compute services: All-purpose, Job, SQL Warehouses, Serverless · Cost models	

PHASE 2 · Weeks 3–5	Data Ingestion & Loading (Domain 2)	21%
Wk 3	COPY INTO command · Incremental loading from ADLS/S3/GCS · Idempotent ingestion · File format options (CSV, JSON, Parquet, Avro)	
Wk 4	Auto Loader deep dive · cloudFiles options · Schema enforcement & evolution · Directory listing vs File notification modes · Checkpointing	
Wk 5	Lakeflow Connect (standard & managed connectors) · JDBC/ODBC/REST ingestion · Semi-structured & nested data (JSON, structs, arrays) · Decision framework	

PHASE 3 · Weeks 6–8	Transformation & Modeling (Domain 3)	22%
Wk 6	PySpark/SQL ETL: bronze → silver tables · Null cleaning · Type casting · Column ops (add, drop, rename, filter, explode arrays)	
Wk 7	All join types (inner, left, broadcast, multi-key, cross, union/unionAll) · Deduplication · Aggregations (count, approx_count_distinct, mean, summary)	
Wk 8	Gold layer objects: materialized views, views, streaming tables, regular tables · Data quality checks · Spark tuning params (shuffle.partitions, autoBroadcast)	

PHASE 4 · Weeks 9–10	Lakeflow Jobs & CI/CD (Domains 4 + 5)	26%
Wk 9	Lakeflow Jobs: notebook, SQL, dashboard, pipeline tasks · DAG dependencies · Retries, conditional branching, loops · Triggers (scheduled, file arrival, table update)	
Wk 10	Databricks Git Folders: branches, commits, PRs · Declarative Automation Bundles (DABs) · Environment promotion (dev/test/prod) · Databricks CLI	

PHASE 5 · Week 11	Monitoring, Optimization & Security (Domains 6 + 7)	25%
Wk 11a	Spark UI: stage metrics, identifying data skew, shuffling, disk spill · Liquid Clustering · Predictive Optimization · OOM diagnosis · Cluster startup issues	
Wk 11b	Unity Catalog: managed vs external tables · GRANT/REVOKE/DENY · Column masking · Row-level security · ABAC policies	

PHASE 6 · Week 12	Final Prep & Mock Exams	Exam
Day 1–3	Take 2 full-length mock exams · Identify weak domains · Review official sample questions from exam guide	
Day 4–5	Deep review of weakest 2 domains · Re-read Delta Lake & Unity Catalog docs	
Day 6–7	Light review · Cheat-sheet skim · Rest · Schedule and take the exam ■	

Databricks Certified

Domain 2: Data Ingestion (21%)	Auto Loader (cloudFiles), COPY INTO, Lakeflow Connect, schema evolution, directory listing vs file notification, JSON/nested data, JDBC/REST
Domain 3: Transformation (22%)	Bronze→Silver→Gold, joins (incl. broadcast), explode arrays, deduplication, aggregates, materialized views vs streaming tables, data quality checks
Domain 7: Governance (15%)	Managed vs external tables, GRANT/REVOKE/DENY, principals, column masking, row-level security, ABAC policies
Domain 4: Lakeflow Jobs (16%)	DAG tasks, dependencies, retries, conditional branching, time vs data-driven triggers, notebook/SQL/dashboard/pipeline task types
Domain 5: CI/CD (10%)	Git Folders, branches, PRs, Declarative Automation Bundles (DABs), environment overrides, Databricks CLI
Domain 6: Optimization (10%)	Spark UI stage metrics, data skew, shuffling, disk spill, Liquid Clustering, Predictive Optimization, OOM, cluster startup
Domain 1: Platform (6%)	Lakehouse architecture, Delta Lake (ACID, time travel), Unity Catalog basics, compute types (all-purpose, job, SQL warehouse, serverless)

■ Recommended Resources & Exam Tips

Official Training (Databricks Academy)	Exam-Day Tips
• **ILT:** Data Engineering with Databricks • Data Ingestion with Lakeflow Connect • Deploy Workloads with Lakeflow Jobs • DevOps Essentials for Data Engineering • Data Interoperability with Unity Catalog • Build Pipelines with Lakeflow Spark Declarative Pipelines • Get Started with Data Governance	✓ Focus on Transformation + Ingestion (43%) ✓ Master Auto Loader options & behavior ✓ Know all join types incl. broadcast hint ✓ Memorize Bronze/Silver/Gold patterns ✓ Understand Spark UI skew indicators ✓ Practice GRANT/REVOKE syntax ✓ Use 90 min wisely: ~2 min/question ✓ Flag uncertain Qs and return later

■ What's New vs v1.0 (Developer Associate)

Removed / De-emphasized	Added / New Focus
• Pure Spark architecture deep-dive (DAG, executors) • Structured Streaming internals • RDD operations • Generic DataFrame API • Detailed Catalyst optimizer	• **Lakeflow Connect** (managed connectors) • **Lakeflow Jobs** orchestration (16%) • **CI/CD with DABs** & Git Folders (10%) • **Unity Catalog** ABAC, masking, RLS • **Liquid Clustering** & Predictive Optimization